

ROBOTIC PHYSICS VIOLATION AUDITS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study foundation-model evaluation for embodied agents. The motivating bet is: Audit robot policies by the physics violations they silently tolerate. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For foundation-model evaluation for embodied agents, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b; 2025a; 2026c;d; 2025b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.535	0.465
observed_only_wam	hidden_mode_shift	0.177	0.823
augmentation_as_data_wam	nominal	0.645	0.355
augmentation_as_data_wam	hidden_mode_shift	0.247	0.752
scalar_uncertainty_wam	nominal	0.670	0.330
scalar_uncertainty_wam	hidden_mode_shift	0.302	0.698
proposed_branch_mechanism	nominal	0.780	0.220
proposed_branch_mechanism	hidden_mode_shift	0.520	0.480

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Robotic Physics Violation Audits argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Agentsafe: Benchmarking the safety of embodied agents on hazardous instructions, 2025a. <http://arxiv.org/abs/2506.14697v3>.
- Transforming monolithic foundation models into embodied multi-agent architectures for human-robot collaboration, 2025b. <http://arxiv.org/abs/2512.00797v1>.
- Hy-embodied-0.5: Embodied foundation models for real-world agents, 2026a. <http://arxiv.org/abs/2604.07430v1>.
- Physics-informed embodied intelligence in the foundation model era: Advancing robot manipulation for smart manufacturing, 2026b. <https://doi.org/10.1016/j.aei.2026.104370>.
- Omnieeg-bench: A standardized evaluation benchmark for eeg foundation models, 2026c. <http://arxiv.org/abs/2606.00815v1>.
- Worldbench: Disambiguating physics for diagnostic evaluation of world models, 2026d. <http://arxiv.org/abs/2601.21282v1>.